

quanteda

October 11, 2014

Type Package

Title Quantitative Analysis of Textual Data

Version 0.5.6.000

Date 2014-10-11

Maintainer Ken Benoit <kbenoit@lse.ac.uk> and Paul Nulty <p.nulty@lse.ac.uk>

Description A package for the management and quantitative analysis of textual data with R. quanteda makes it easy to manage texts in the form of a corpus, defined as a collection of texts that includes document-level variables specific to each text, as well as meta-data for documents and for the collection as a whole. quanteda includes tools to make it easy and fast to manipulate the texts in a corpus, for instance by tokenizing them, with or without stopwords or stemming, or to segment them by sentence or paragraph units. quanteda implements bootstrapping methods for texts that makes it easy to resample texts from pre-defined units, to facilitate computation of confidence intervals on textual statistics using techniques of non-parametric bootstrapping, but applied to the original texts as data. quanteda includes a suite of sophisticated tools to extract features of the texts into a quantitative matrix, where these features can be defined according to a dictionary or thesaurus, including the declaration of collocations to be treated as single features. Once converted into a quantitative matrix (known as a ``dfm" for document-feature matrix), the textual feature can be analyzed using quantitative methods for describing, comparing, or scaling texts, or used to train machine learning methods for class prediction.

Encoding UTF-8

License GPL-3

Depends R (>= 3.0), data.table

Imports data.table, SnowballC, wordcloud, slam, tm, proxy

Suggests

quantedaData, austin, entropy, jsonlite, openNLP, RJSONIO, RCurl, twitteR, XML, lda, topicmodels, tcltk2, knitr

URL <http://github.com/kbenoit/quanteda>

BugReports <https://github.com/kbenoit/quanteda/issues>

LazyData TRUE

VignetteBuilder knitr

R topics documented:

bigrams	3
changeunits	4
clean	4
collocations	5
corpus	7
countSyllables	10
describeTexts	10
dfm	11
dfm2ldaformat	13
dfm2tmformat	13
directory	14
docnames	14
docvars	15
encoding	16
features.dfm	17
flatten.dictionary	17
getRootFileNames	18
getTextDir	19
getTextDirGui	19
getTextFiles	20
getTweets	21
inaugCorpus	21
kwic	22
language	23
metacorpus	23
metadoc	24
ndoc	25
ngrams	25
nresample	26
plot.dfm	27
print.dfm	27
quanteda	28
readLIWCdict	28
readWStatDict	29
resample	30
segmentSentence	31
settings	32
settingsInitialize	33
similarity	34
sort.dfm	35
statLexdiv	36
stopwords	38
stopwordsGet	39
stopwordsRemove	39
subset.corpus	40
summary.corpus	41

<i>bigrams</i>	3
syllableCounts	41
texts	42
tf	42
tfidf	43
tokenize	44
topfeatures	45
trimdfm	45
uk2010immig	46
wordstem	47
zipfiles	47
Index	49

bigrams	<i>Create bigrams</i>
---------	-----------------------

Description

Create bigrams

Usage

```
bigrams(text, window = 1, concatenator = "_", include.unigrams = FALSE,
...)
```

Arguments

text	character vector containing the texts from which bigrams will be constructed
window	how many words to be counted for adjacency. Default is 1 for only immediately neighbouring words. This is only available for bigrams, not for ngrams.
concatenator	character for combining words, default is _ (underscore) character
include.unigrams	if TRUE, return unigrams as well
...	provides additional arguments passed to tokenize

Value

a character vector of bigrams

Author(s)

Ken Benoit and Kohei Watanabe

Examples

```
bigrams("The quick brown fox jumped over the lazy dog.")
bigrams(c("The quick brown fox", "jumped over the lazy dog. "))
bigrams(c("The quick brown fox", "jumped over the lazy dog."), window=2)
```

changeunits

change the document units of a corpus

Description

For a corpus, recast the documents down or up a level of aggregation. "Down" would mean going from documents to sentences, for instance. "Up" means from sentences back to documents. This makes it easy to reshape a corpus from a a collection of documents into a collection of sentences, for instance.

Usage

```
changeunits(corp, to = c("sentences", "paragraphs", "documents"), ...)
```

Arguments

corp	corpus whose document units will be reshaped
to	new documents units for the corpus to be recast in
...	passes additional arguments to segment

Examples

```
# simple example
mycorpus <- corpus(c(textone="This is a sentence. Another sentence. Yet another.",
                     texttwo="Premiere phrase. Deuxieme phrase."),
                  docvars=list(country=c("UK", "USA"), year=c(1990, 2000)),
                  notes="This is a simple example to show how changeunits() works.")
language(mycorpus) <- c("english", "french")
summary(mycorpus)
summary(changeunits(mycorpus, to="sentences"), showmeta=TRUE)

# example with inaugural corpus speeches
mycorpus2 <- subset(inaugCorpus, Year>2004)
mycorpus2
paragCorpus <- changeunits(mycorpus2, to="paragraphs")
paragCorpus
summary(paragCorpus, 100, showmeta=TRUE)
## Note that Bush 2005 is recorded as a single paragraph because that text used a single
## \n to mark the end of a paragraph.
```

clean

simple cleaning of text before processing

Description

clean removes punctuation and digits from text, using the regex character classes for punctuation and digits. clean uses the standard R function tolower to convert the text to lower case. Each of these steps is optional, but switched on by default, so for example, to remove punctuation and convert to lower, but keep digits, the command would be: clean(mytexts, removeDigits=FALSE)

Usage

```
clean(x, ...)

## S3 method for class character
clean(x, removeDigits = TRUE, removePunct = TRUE,
      lower = TRUE, additional = NULL, twitter = FALSE, ...)

## S3 method for class corpus
clean(x, removeDigits = TRUE, removePunct = TRUE,
      lower = TRUE, additional = NULL, twitter = FALSE, ...)
```

Arguments

<code>x</code>	The object to be cleaned. Can be either a character vector or a corpus object. If <code>x</code> is a corpus, <code>clean</code> returns a copy of the <code>x</code> with the texts cleaned.
<code>...</code>	additional parameters
<code>removeDigits</code>	remove numbers if TRUE
<code>removePunct</code>	remove punctuation if TRUE
<code>lower</code>	convert text to lower case TRUE
<code>additional</code>	additional characters to remove (regular expression)
<code>twitter</code>	if TRUE, do not remove @ or #

Value

A character vector equal in length to the original texts, after cleaning.

Examples

```
clean("This is 1 sentence with 1.9 numbers in it, and one comma.", removeDigits=FALSE)
clean("This is 1 sentence with 1.9 numbers in it, and one comma.", lower=FALSE)
clean("The \"economy\" is [worth] approx. $1.2 trillion.")
clean("The \"economy\" is {worth} approx. $1.2 trillion.", additional="[$]")
clean("We are his Beliebers, and him is #ourjustin @justinbieber we love u", twitter=TRUE)
clean("Collocations can be represented as inheritance_tax using the _ character.")

# for a vector of texts
clean(c("This is 1 sentence with 1.9 numbers in it, and one comma.",
        "$1.2 billion was spent on text analysis in 2014."))
```

collocations

Detect collocations from text

Description

Detects collocations (currently, bigrams) from texts or a corpus, returning a data.frame of collocations and their scores, sorted by the likelihood ratio G^2 and Pearson's χ^2 .

Usage

```
collocations(x, ...)

## S3 method for class character
collocations(x, method = c("lr", "chi2", "pmi", "dice",
  "all"), n = 2, top = NULL, ...)

## S3 method for class corpus
collocations(x, method = c("lr", "chi2", "pmi", "dice",
  "all"), n = 2, top = NULL, ...)
```

Arguments

x	a text, a character vector of texts, or a corpus
...	additional parameters
method	association measure for detecting collocations. Available measures for bigrams are: "lr" The likelihood ratio statistic G^2 , computed as: $2 * \sum_i \sum_j (n_{ij} * \log \frac{n_{ij}}{m_{ij}})$ "chi2" Pearson's χ^2 statistic, computed as: $\sum_i \sum_j \frac{(n_{ij} - m_{ij})^2}{m_{ij}}$ "pmi" point-wise mutual information score, computed as $\log n_{11}/m_{11}$ "dice" the Dice coefficient, computed as $n_{11}/n_{1.} + n_{.1}$ "all" returns all of the above
n	length of the collocation. Only bigrams (n=2) implemented so far.
top	the number of collocations to return, sorted in descending order of the requested statistic, or G^2 if none is specified.

Value

A data.frame of collocations, their frequencies, and the computed association measure.

Author(s)

Kenneth Benoit

References

Add some.

See Also

bigrams, trigrams

Examples

```
collocations(inaugTexts, top=10)
collocations(inaugCorpus, top=10, method="chi2")
```

corpus

*Constructor for corpus objects***Description**

Creates a corpus from a document source. The current available document sources are:

- a character vector (as in R class `char`) of texts;
- a directory of text files, using [directory](#);
- a directory constructed from a zip file consisting of text files, using [zipfiles](#); and
- a **tm** [VCorpus](#) class corpus object, meaning that anything you can use to create a **tm** corpus, including all of the tm plugins plus the built-in functions of tm for importing pdf, Word, and XML documents, can be used to create a quanteda [corpus](#).

Corpus-level meta-data can be specified at creation, containing (for example) citation information and notes, as can document-level variables and document-level meta-data.

Usage

```
corpus(x, ...)
```

```
## S3 method for class directory
corpus(x, enc = NULL, docnames = NULL,
       docvarsfrom = c("none", "filenames", "headers"), docvarnames = NULL,
       sep = "_", pattern = "\\..txt$", source = NULL, notes = NULL,
       citation = NULL, ...)
```

```
## S3 method for class twitter
corpus(x, enc = NULL, notes = NULL, citation = NULL,
       ...)
```

```
## S3 method for class VCorpus
corpus(x, enc = NULL, notes = NULL, citation = NULL,
       ...)
```

```
## S3 method for class character
corpus(x, enc = NULL, docnames = NULL, docvars = NULL,
       source = NULL, notes = NULL, citation = NULL, ...)
```

```
is.corpus(x)
```

```
## S3 method for class corpus
c1 + c2
```

Arguments

<code>x</code>	A source of texts to form the documents in the corpus. This can be a filepath to a directory containing text documents (see directory), or a character vector of texts.
<code>...</code>	additional arguments

enc	A string specifying the input encoding for texts in the corpus. Must be a valid entry in <code>iconvlist()</code> , since the code in <code>corpus.character</code> will convert this to UTF-8 using <code>iconv</code> . Currently only one input encoding can be specified for a collection of input texts, meaning that you should not mix input text encoding types in a single corpus call.
docnames	Names to be assigned to the texts, defaults to the names of the character vector (if any), otherwise assigns "text1", "text2", etc.
docvarsfrom	Argument to specify where docvars are to be taken, from parsing the filenames separated by <code>sep</code> or from meta-data embedded in the text file header (headers).
docvarnames	Character vector of variable names for docvars
sep	Separator if <code>docvars</code> names are taken from the filenames.
pattern	filename extension - set to "*" if all files are desired. This is a regular expression .
source	A string specifying the source of the texts, used for referencing.
notes	A string containing notes about who created the text, warnings, To Dos, etc.
citation	Information on how to cite the corpus.
docvars	A data frame of attributes that is associated with each text.
c1	corpus one to be added
c2	corpus two to be added

Details

The `+` operator for a corpus object will combine two corpus objects, resolving any non-matching `docvars` or `metadoc` fields by making them into NA values for the corpus lacking that field. Corpus-level meta data is concatenated, except for source and notes, which are stamped with information pertaining to the creation of the new joined corpus.

There are some issues that need to be addressed in future revisions of `quanteda` concerning the use of factors to store document variables and meta-data. Currently most or all of these are not recorded as factors, because we use `stringsAsFactors=FALSE` in the `data.frame` calls that are used to create and store the document-level information, because the texts should always be stored as character vectors and never as factors.

Value

A corpus class object containing the original texts, document-level variables, document-level meta-data, corpus-level metadata, and default settings for subsequent processing of the corpus. A corpus consists of a list of elements described below, although these should only be accessed through accessor and replacement functions, not directly (since the internals may be subject to change). The structure of a corpus classed list object is:

<code>\$documents</code>	A data frame containing the document level information, consisting of <code>texts</code> , user-named <code>docvars</code> variables describing attributes of the documents, and <code>metadoc</code> document-level metadata whose names begin with an underscore character, such as <code>_language</code> .
<code>\$metadata</code>	A named list set of corpus-level meta-data, including source and created (both generated automatically unless assigned), notes, and citation.
<code>\$settings</code>	Settings for the corpus which record options that govern the subsequent processing of the corpus when it is converted into a document-feature matrix (<code>dfm</code>). See settings .

`$tokens` An indexed list of tokens and types tabulated by document, including information on positions. Not yet fully implemented.

`is.corpus` returns TRUE if the object is a corpus

Note

When `x` is a [VCorpus](#) object, the fixed metadata fields from that object are imported as document-level metadata. Currently no corpus-level metadata is imported, but we will add that soon.

See Also

[docvars](#), [metadoc](#), [metacorporus](#), [language](#), [encoding](#), [settings](#), [texts](#)

Examples

```
## Not run:
# import texts from a directory of files
summary(corpus(directory("~/Dropbox/QUANTESS/corpora/ukManRenamed"),
  enc="UTF-8", docvarsfrom="filenames",
  source="Kens UK manifesto archive",
  docvarnames=c("Country", "Level", "Year", "language")), 5))
summary(corpus(directory("~/Dropbox/QUANTESS/corpora/ukManRenamed"),
  enc="UTF-8", docvarsfrom="filenames",
  source="Kens UK manifesto archive",
  docvarnames=c("Country", "Level", "Year", "language", "Party")), 5))

# choose a directory using a GUI
corpus(directory())

# from a zip file on the web
myzipcorp <- corpus(zipfiles("http://kenbenoit.net/files/EUcoalsubsidies.zip"),
  notes="From some EP debate about coal mine subsidies")
docvars(myzipcorp, speakername=docnames(myzipcorp))
summary(myzipcorp)

## End(Not run)
#
## import a tm VCorpus
if (require(tm)) {
  data(crude) # load in a tm example VCorpus
  mytmCorpus <- corpus(crude)
  summary(mytmCorpus, showmeta=TRUE)
}
#
# create a corpus from texts
corpus(inaugTexts)

# create a corpus from texts and assign meta-data and document variables
uk2010immigCorpus <- corpus(uk2010immig,
  docvars=data.frame(party=names(uk2010immig)),
  enc="UTF-8")
```

countSyllables	Returns a count of the number of syllables in the input
----------------	---

Description

This function takes a text and returns a count of the number of syllables it contains. For British English words, the syllable count is exact and looked up from the CMU pronunciation dictionary. For any word not in the dictionary the syllable count is estimated by counting vowel clusters.

Usage

```
countSyllables(sourceText)
```

Arguments

sourceText	Character vector of texts whose syllables will be counted
------------	---

Details

This only works for English.

Value

numeric Named vector of counts of the number of syllables for each element of sourceText. When a word is not available in the lookup table, its syllables are estimated by counting the number of (English) vowels in the word.

Examples

```
countSyllables("This is an example sentence.")
myTexts <- c("Text one.", "Superduper text number two.", "One more for the road.")
names(myTexts) <- paste("myText", 1:3, sep="")
countSyllables(myTexts)
```

describeTexts	print a summary of texts
---------------	--------------------------

Description

Prints to the console a description of the texts, including number of types, tokens, and sentences

Usage

```
describeTexts(txts, verbose = TRUE)
```

Arguments

txts	The texts to be described
verbose	Default is TRUE. Set to false to suppress output messages

Examples

```
describeTexts(c("testing this text", "and this one"))
describeTexts(uk2010immig)
```

dfm

Create a document-feature matrix from a corpus object

Description

Returns a document by feature matrix with additional meta-information (settings, identification of training texts for supervised models, resampling information, etc.) that is useful in other quanteda functions. A typical usage would be to produce a word-frequency matrix where the cells are counts of words by document, but the definition of "features" is entirely general.

Usage

```
dfm(x, ...)

## S3 method for class corpus
dfm(x, feature = c("word"), stem = FALSE,
    stopwords = NULL, bigram = FALSE, groups = NULL, verbose = TRUE,
    dictionary = NULL, dictionary_regex = FALSE, bootstrap = FALSE,
    addto = NULL, ...)

## S3 method for class character
dfm(x, feature = c("word"), stem = FALSE,
    stopwords = NULL, bigram = FALSE, verbose = TRUE, dictionary = NULL,
    dictionary_regex = FALSE, addto = NULL, ...)

is.dfm(x)

as.dfm(x)
```

Arguments

x	Corpus or character vector from which to generate the document-feature matrix
...	additional arguments passed to clean
feature	Feature to count (e.g. words)
stem	Stem the words
stopwords	A character vector of stopwords that will be removed from the text when constructing the dfm . If NULL (default) then no stopwords will be applied. If "TRUE" then it currently defaults to stopwords .
bigram	include bigrams as well as unigram features, if TRUE
groups	Grouping variable for aggregating documents
verbose	Get info to screen on the progress
dictionary	A list of character vector dictionary entries, including regular expressions (see examples)

dictionary_regex	TRUE means the dictionary is already in regular expression format, otherwise it will be converted from "wildcard" format
bootstrap	if TRUE, compute multiple dfm's from resampled texts in the corpus. Requires a resampled corpus. See resample .
addto	NULL by default, but if an existing dfm object is specified, then the new dfm will be added to the one named. If both dfm 's are built from dictionaries, the combined dfm will have its Non_Dictionary total adjusted.

Details

`is.dfm` returns TRUE if and only if its argument is a [dfm](#).

`as.dfm` coerces a matrix to a dfm

Value

A specially classed matrix object with row names equal to the document names and column names equal to the feature labels. Additional information is attached to this object as [attributes](#), such as [settings](#).

Author(s)

Kenneth Benoit

Examples

```
wfm <- dfm(inaugCorpus)

## by president, after 1960
wfmByPresfrom1900 <- dfm(subset(inaugCorpus, Year>1900), groups="President")
docnames(wfmByPresfrom1900)

## with dictionaries
mycorpus <- subset(inaugCorpus, Year>1900)
mydict <- list(christmas=c("Christmas", "Santa", "holiday"),
               opposition=c("Opposition", "reject", "notincorpus"),
               taxing="taxing",
               taxation="taxation",
               taxregex="tax*")
dictDfm <- dfm(mycorpus, dictionary=mydict)
dictDfm

## removing stopwords
testText <- "The quick brown fox named Seamus jumps over the lazy dog also named Seamus, with
the newspaper from a a boy named Seamus, in his mouth."
testCorpus <- corpus(testText)
settings(testCorpus, "stopwords")
dfm(testCorpus, stopwords=TRUE)
```

dfm2ldaformat	Convert a dfm into the format needed by lda
---------------	--

Description

Convert a quanteda [dfm](#) object into the indexed format required by the topic modelling package **lda**.

Usage

```
dfm2ldaformat(d)
```

Arguments

d A [dfm](#) object

Value

A list with components "documents" and "vocab" as needed by [lda.collapsed.gibbs.sampler](#)

Examples

```
mycorpus <- subset(inaugCorpus, Year>1970)
d <- dfm(mycorpus, stopwords=TRUE)
d <- trimdfm(d, minCount=5, minDoc=3)
td <- dfm2ldaformat(d)
if (require(lda)) {
  tmodel.lda <- lda.collapsed.gibbs.sampler(documents=td$documents,
                                           K=10,
                                           vocab=td$vocab,
                                           num.iterations=50, alpha=0.1, eta=0.1)
  top.topic.words(tmodel.lda$topics, 10, by.score=TRUE) # top five words in each topic
}
```

dfm2tmformat	Convert a dfm into a tm DocumentTermMatrix
--------------	---

Description

tm represents sparse document-feature matrixes in the [simple triplet matrix](#) format of the package **slam**. This function converts a dfm into a [DocumentTermMatrix](#), enabling a dfm to be used with other packages that expect this format, such as **topicmodels**.

Usage

```
dfm2tmformat(d, weighting = weightTf)
```

Arguments

d A [dfm](#) object

weighting weight function arguments passed to `as.TermDocumentMatrix`, defaults to term frequency (see [as.DocumentTermMatrix](#) for a list of options, such as tf-idf).

Value

A simple triplet matrix of class `as.DocumentTermMatrix`

Examples

```
mycorpus <- subset(inaugCorpus, Year>1970)
d <- trimdfm(dfm(mycorpus), minCount=5, minDoc=3)
dim(d)
td <- dfm2tmformat(d)
length(td$v)
if (require(topicmodels)) (tmodel.lda <- LDA(td, control = list(alpha = 0.1), k = 5))
```

directory

Function to declare a connection to a directory (containing files)

Description

Function to declare a connection to a directory, although unlike `file` it does not require closing. If the directory does not exist, the function will return an error.

Usage

```
directory(path = NULL)
```

Arguments

path	String describing the full path of the directory or NULL to use a GUI to choose a directory from disk
------	---

Examples

```
## Not run:
# name a directory of files
mydir <- directory("~/Dropbox/QUANTESS/corpora/ukManRenamed")
corpus(mydir)

# choose a directory using a GUI
corpus(directory())
## End(Not run)
```

docnames

get or set document names

Description

Extract the document names from a corpus or a document-feature matrix. Document names are the rownames of the documents data.frame in a corpus, or the rownames of the `dfm` object for a dfm. of the `dfm` object.

`docnames` queries the document names of a corpus or a dfm

`docnames <-` assigns new values to the document names of a corpus. (Does not work for dfm objects, whose document names are fixed.)

Usage

```
docnames(x)

## S3 method for class corpus
docnames(x)

docnames(x) <- value

## S3 method for class dfm
docnames(x)
```

Arguments

x the object with docnames
 value a character vector of the same length as x

Value

docnames returns a character vector of the document names
 docnames<- assigns a character vector of the document names in a corpus

Examples

```
# query the document names of the inaugural speech corpus
docnames(inaugCorpus) <- paste("Speech", 1:ndoc(inaugCorpus), sep="")

# reassign the document names of the inaugural speech corpus
docnames(inaugCorpus) <- paste("Speech", 1:ndoc(inaugCorpus), sep="")
#
# query the document names of a dfm
docnames(dfm(inaugTexts[1:5]))
```

docvars	<i>get or set for document-level variables</i>
---------	--

Description

Get or set variables for the documents in a corpus

Usage

```
docvars(x, field = NULL)

docvars(x, field = NULL) <- value
```

Arguments

x corpus whose document-level variables will be read or set
 field string containing the document-level variable name
 value the new values of the document-level variable

Value

`docvars` returns a data.frame of the document-level variables

`docvars<-` assigns value to the named field

Examples

```
head(docvars(inaugCorpus))
docvars(inaugCorpus, "President") <- paste("prez", 1:ndoc(inaugCorpus), sep="")
head(docvars(inaugCorpus))
```

encoding	<i>get the encoding of documents in a corpus</i>
----------	--

Description

Get or set the `_encoding` document-level metadata field in a corpus.

Usage

```
encoding(x)
```

```
encoding(x) <- value
```

Arguments

`x` a corpus object

`value` a character vector or scalar representing the new value of the encoding (see Note)

Details

This function modifies the `_encoding` value set by [metadoc](#). It is a wrapper for `metadoc(corp, "encoding")`.

Note

This function differs from R's built-in [Encoding](#) function, which only allows the four values of "latin1", "UTF-8", "bytes", and "unknown" (and which assigns "unknown" to any text that contains only ASCII characters). Legal values for encodings must be from [iconvlist](#). Note that encoding does not convert or set encodings, it simply records a user declaration of a valid encoding. (We hope to implement checking and conversion later.)

features.dfm	<i>extract the feature labels from a dfm</i>
--------------	--

Description

Extract the features from a document-feature matrix, which are stored as the column names of the [dfm](#) object.

Usage

```
## S3 method for class dfm
features(x)
```

Arguments

x the object (dfm) whose features will be extracted

Value

Character vector of the features

Examples

```
features(dfm(inaugTexts))[1:50] # first 50 features (alphabetically sorted)
```

flatten.dictionary	<i>Flatten a hierarchical dictionary into a list of character vectors</i>
--------------------	---

Description

Converts a hierarchical dictionary (a named list of named lists, ending in character vectors at the lowest level) into a flat list of character vectors. Works like `unlist(dictionary, recursive=TRUE)` except that the recursion does not go to the bottom level.

Usage

```
flatten.dictionary(elms, parent = "", dict = list())
```

Arguments

elms	list to be flattened
parent	parent list name, gets built up through recursion in the same way that <code>unlist(dictionary, recursive=TRUE)</code> works
dict	the bottom list of dictionary entries ("synonyms") passed up from recursive calls

Details

Called by `dfm()`

Value

A dictionary flattened down one level further than the one passed

Author(s)

Kohei Watanabe

Examples

```
dictPopulismEN <-
  list(populism=c("elit*", "consensus*", "undemocratic*", "referend*",
    "corrupt*", "propagand", "politici*", "*deceit*",
    "*deceiv*", "*betray*", "shame*", "scandal*", "truth*",
    "dishonest*", "establishm*", "ruling*"))
flatten.dictionary(dictPopulismEN)

hdict <- list(level1a = list(level1a1 = c("l1a11", "l1a12"),
  level1a2 = c("l1a21", "l1a22")),
  level1b = list(level1b1 = c("l1b11", "l1b12"),
    level1b2 = c("l1b21", "l1b22", "l1b23")),
  level1c = list(level1c1a = list(level1c1a1 = c("lowest1", "lowest2")),
    level1c1b = list(level1c1b1 = c("lowestalone"))))
flatten.dictionary(hdict)
```

getRootFileNames

Truncate absolute filepaths to root filenames

Description

This function takes an absolute filepath and returns just the document name

Usage

```
getRootFileNames(longFilenames)
```

Arguments

longFilenames Absolute filenames including a full path with directory

Value

character vector of filenames without directory path

Author(s)

Paul Nulty

Examples

```
## Not run:
getRootFileNames(/home/paul/documents/libdem09.txt)

## End(Not run)
```

getTextDir	<i>loads all text files from a given directory</i>
------------	--

Description

given a directory name, get a list of all files in that directory and load them into a character vector using getTextFiles

Usage

```
getTextDir(dirname, pattern = "\\*.txt$", enc = "unknown")
```

Arguments

dirname	A directory path
pattern	a regular expression pattern match for the input file names
enc	a value for encoding that is a legal value for Encoding

Value

character vector of texts read from disk

Author(s)

Paul Nulty

Examples

```
## Not run:  
getTextDir(/home/paul/documents/)  
  
## End(Not run)
```

getTextDirGui	<i>provides a gui interface to choose a gui to load texts from</i>
---------------	--

Description

launches a GUI to allow the user to choose a directory from which to load all files.

Usage

```
getTextDirGui()
```

Value

character vector of texts read from disk

Author(s)

Paul Nulty

Examples

```
## Not run:
getTextFiles(/home/paul/documents/libdem09.txt)

## End(Not run)
```

getTextFiles	<i>load text files from disk into a vector of character vectors points to files, reads them into a character vector of the texts with optional names, default being filenames returns a named vector of complete, unedited texts</i>
--------------	--

Description

load text files from disk into a vector of character vectors points to files, reads them into a character vector of the texts with optional names, default being filenames returns a named vector of complete, unedited texts

Usage

```
getTextFiles(filenamees, textnames = NULL, enc = "unknown",
  verbose = FALSE)
```

Arguments

filenamees	a vector of paths to text files
textnames	names to assign to the texts
enc	a value for encoding that is a legal value for Encoding
verbose	If TRUE, print out names of files being read. Default is FALSE

Value

character vector of texts read from disk

Author(s)

Paul Nulty

Examples

```
## Not run:
getTextFiles(/home/paul/documents/libdem09.txt)

## End(Not run)
```

getTweets	<i>Function to declare a twitter search</i>
-----------	---

Description

Function to declare a connection to a twitter search

Usage

```
getTweets(query, numResults = 50, key, cons_secret, token, access_secret)
```

Arguments

query	String describing the search query terms
numResults	Number of tweets to return. Maximum of approximately 1500
key	Key for twitter API authentication
cons_secret	for twitter API authentication
token	String for twitter API authentication
access_secret	for twitter API authentication

Value

The search results marked as a 'twitter' object for use by corpus.twitter()

inaugCorpus	<i>A corpus of US presidential inaugural addresses from 1789-2013</i>
-------------	---

Description

inaugCorpus is the [quanteda](#) corpus object of US presidents' inaugural addresses since 1789. Document variables contain the year of the address and the last name of the president.

inaugTexts is the character vector of US presidential inauguration speeches

References

<https://archive.org/details/Inaugural-Address-Corpus-1789-2009> and <http://www.presidency.ucsbs.edu/inaugurals.php>.

Examples

```
# some operations on the inaugural corpus
data(inaugCorpus)
summary(inaugCorpus)
head(docvars(inaugCorpus), 10)
# working with the character vector only
data(inaugTexts)
str(inaugTexts)
head(docvars(inaugCorpus), 10)
mycorpus <- corpus(inaugTexts)
```

kwic

List key words in context from a text or a corpus of texts.

Description

For a text or a collection of texts (in a quanteda corpus object), return a list of a keyword supplied by the user in its immediate context, identifying the source text and the word index number within the source text. (Not the line number, since the text may or may not be segmented using end-of-line delimiters.)

Usage

```
kwic(x, word, window = 5, regex = TRUE)
```

```
## S3 method for class character
```

```
kwic(x, word, window = 5, regex = TRUE)
```

```
## S3 method for class corpus
```

```
kwic(x, word, window = 5, regex = TRUE)
```

Arguments

x	A text character scalar or a quanteda corpus. (Currently does not support character vectors.)
word	A keyword chosen by the user.
window	The number of context words to be displayed around the keyword.
regex	If TRUE (default), then "word" is a regular expression, otherwise only match the whole word. Note that if regex=TRUE and no special regular expression characters are used in the search query, then the concordance will include all words in which the search term appears, and not just when it appears as an entire word. (For instance, searching for the word "key" will also return "whiskey".)

Value

A data frame with the context before (preword), the keyword in its original format (word, preserving case and attached punctuation), and the context after (postword). The rows of the dataframe will be named with the word index position, or the text name and the index position for a corpus object.

Author(s)

Kenneth Benoit and Paul Nulty

Examples

```
kwic(inaugTexts, "terror")
kwic(inaugTexts, "terror", regex=FALSE) # returns only whole word, without trailing punctuation
```

language	<i>get or set the language of corpus documents</i>
----------	--

Description

Get or set the `_language` document-level metadata field in a corpus.

Usage

```
language(corp)
```

```
language(corp) <- value
```

Arguments

corp	a corpus object
value	the new value for the language meta-data field, a string or character vector equal in length to <code>ndoc(corp)</code>

Details

This function modifies the `_language` value set by [metadoc](#). It is a wrapper for `metadoc(corp, "language")`.

metacorporus	<i>get or set corpus metadata</i>
--------------	-----------------------------------

Description

Get or set the corpus-level metadata in a quanteda corpus object.

Usage

```
metacorporus(corp, field = NULL)
```

```
metacorporus(corp, field) <- value
```

Arguments

corp	A quanteda corpus object
field	Metadata field name(s). If NULL (default), return all metadata names.
value	new value of the corpus metadata field

Value

For `metacorporus`, a list of the metadata fields in the corpus. If a list is not what you wanted, you can wrap the results in [unlist](#), but this will remove any metadata field that is set to NULL.

For `metacorporus <-`, the corpus with the updated metadata.

Examples

```
metacorp(inaugCorpus)
metacorp(inaugCorpus, "source")
metacorp(inaugCorpus, "citation") <- "Presidential Speeches Online Project (2014)."
```

```
metacorp(inaugCorpus, "citation")
```

metadoc	<i>get or set document-level meta-data</i>
---------	--

Description

Get or set the document-level meta-data, including reserved fields for language and corpus.

Usage

```
metadoc(corp, field = NULL)

metadoc(corp, field = NULL) <- value
```

Arguments

corp	A quanteda corpus object
field	string containing the name of the metadata field(s) to be queried or set
value	the new value of the new meta-data field

Value

For texts, a character vector of the texts in the corpus.

For texts <-, the corpus with the updated texts.

Note

Document-level meta-data names are preceded by an underscore character, such as `_encoding`, but when named in in the `field` argument, do *not* need the underscore character.

Examples

```
mycorp <- subset(inaugCorpus, Year>1990)
summary(mycorp, showmeta=TRUE)
metadoc(mycorp, "encoding") <- "UTF-8"
metadoc(mycorp)
metadoc(mycorp, "language") <- "english"
summary(mycorp, showmeta=TRUE)
```

ndoc	<i>get the number of documents</i>
------	------------------------------------

Description

Returns the number of documents in a corpus objects

Usage

```
ndoc(x)

## S3 method for class corpus
ndoc(x)

## S3 method for class dfm
ndoc(x, ...)
```

Arguments

x	a corpus or dfm object
...	additional parameters

Value

an integer (count) of the number of documents in the corpus or dfm

Examples

```
ndoc(inaugCorpus)
ndoc(dfm(inaugCorpus))
```

ngrams	<i>Create ngrams</i>
--------	----------------------

Description

Create a set of ngrams (words in sequence) from text(s) in a character vector

Usage

```
ngrams(text, n = 2, concatenator = "_", include.all = FALSE, ...)
```

Arguments

text	character vector containing the texts from which ngrams will be extracted
n	the number of tokens to concatenate. Default is 2 for bigrams.
concatenator	character for combining words, default is _ (underscore) character
include.all	if TRUE, add n-1...1 grams to the returned list
...	additional parameters

Details

... provides additional arguments passed to [tokenize](#)

Value

a list of character vectors of ngrams, one list element per text

Author(s)

Ken Benoit, Kohei Watanabe, Paul Nulty

Examples

```

ngrams("The quick brown fox jumped over the lazy dog.", n=2)
identical(ngrams("The quick brown fox jumped over the lazy dog.", n=2),
          bigrams("The quick brown fox jumped over the lazy dog.", n=2))
ngrams("The quick brown fox jumped over the lazy dog.", n=3)
ngrams("The quick brown fox jumped over the lazy dog.", n=3, concatenator="~")
ngrams("The quick brown fox jumped over the lazy dog.", n=3, include.all=TRUE)

```

nresample

get the number of resamples

Description

Get the number of resamples from a corpus or dfm object

Usage

```

nresample(x)

## S3 method for class corpus
nresample(x)

## S3 method for class dfm
nresample(x)

```

Arguments

x corpus object containing the texts to be resampled

Value

an integer as the number of resampled texts

plot.dfm	<i>plot features as a wordcloud</i>
----------	-------------------------------------

Description

The default plot method for a [dfm](#) object. Produces a wordcloud plot for the features of the dfm, weighted by the total frequencies. To produce word cloud plots for specific documents, the only way currently to do this is to produce a dfm only from the documents whose features you want plotted.

Usage

```
## S3 method for class dfm
plot(x, ...)
```

Arguments

x	a dfm object
...	additional parameters passed to wordcloud or to text (and strheight , strwidth)

See Also

[wordcloud](#)

Examples

```
# plot the features (without stopwords) from Obamas two inaugural addresses
mydfm <- dfm(subset(inaugCorpus, President=="Obama"), verbose=FALSE, stopwords=TRUE)
plot(mydfm)

# plot only Lincolns inaugural address
plot(dfm(subset(inaugCorpus, President=="Lincoln"), verbose=FALSE, stopwords=TRUE))

# plot in colors with some additional options passed to wordcloud
plot(mydfm, random.color=TRUE, rot.per=.25, colors=sample(colors()[2:128], 5))
```

print.dfm	<i>print a dfm object</i>
-----------	---------------------------

Description

print method for dfm objects

Usage

```
## S3 method for class dfm
print(x, show.values = FALSE, show.settings = FALSE, ...)
```

Arguments

x	the dfm to be printed
show.values	print the dfm as a matrix or array (if resampled).
show.settings	Print the settings used to create the dfm. See settings .
...	further arguments passed to or from other methods

 quanteda

An R package for the quantitative analysis of textual data.

Description

A set of functions for creating and managing text corpora, extracting features from text corpora, and analyzing those features using quantitative methods.

Author(s)

Ken Benoit and Paul Nulty

 readLIWCdict

Import a LIWC-formatted dictionary

Description

Make a flattened dictionary list object from a LIWC dictionary file.

Usage

```
readLIWCdict(path, maxcats = 10, enc = "")
```

Arguments

path	full pathname of the LIWC-formatted dictionary file (usually a file ending in .dic)
maxcats	the maximum number of categories to read in, set by the maximum number of dictionary categories that a term could belong to. For non-exclusive schemes such as the LIWC, this can be up to 7. Set to 10 by default, which ought to be more than enough.
enc	a valid input encoding for the file to be read, see iconvlist

Value

a dictionary class named list, where each the name of element is a bottom level category in the hierarchical wordstat dictionary. Each element is a list of the dictionary terms corresponding to that level.

Author(s)

Kenneth Benoit

Examples

```
## Not run:
LIWCdict <- readLIWCdict("~/Dropbox/QUANTESS/corpora/LIWC/LIWC2001_English.dic")
## End(Not run)
```

readWStatDict	<i>Import a Wordstat dictionary</i>
---------------	-------------------------------------

Description

Make a flattened list from a hierarchical wordstat dictionary

Usage

```
readWStatDict(path, enc = "", lower = TRUE)
```

Arguments

path	full pathname of the wordstat dictionary file (usually ending in .cat)
enc	a valid input encoding for the file to be read, see iconvlist
lower	if TRUE (default), convert the dictionary entries to lower case

Value

a named list, where each the name of element is a bottom level category in the hierarchical wordstat dictionary. Each element is a list of the dictionary terms corresponding to that level.

Author(s)

Kohei Watanabe, Kenneth Benoit

Examples

```
## Not run:
path <- ~/Dropbox/QUANTESS/corpora/LaverGarry.cat
lgdict <- readWStatDict(path)

## End(Not run)
```

resample*resampling methods for a corpus*

Description

Draw a set of random resamples from a corpus object, at a specified level of resampling, and record additional "resampled texts" as document-level metadata, stored as `_resampleXXX` for the XXXth resample.

Usage

```
resample(x, ...)  
  
## S3 method for class corpus  
resample(x, n = 100, unit = c("sentences", "paragraphs"),  
  ...)  
  
is.resampled(x)  
  
## S3 method for class corpus  
is.resampled(x)  
  
## S3 method for class dfm  
is.resampled(x)
```

Arguments

<code>x</code>	corpus object containing the texts to be resampled
<code>...</code>	additional arguments passed to segment
<code>n</code>	number of resamples to be drawn
<code>unit</code>	resampling unit for drawing the random samples, can be sentences or paragraphs.

Details

`is.resampled` checks a corpus or dfm object and returns TRUE if these contain resampled texts or the results of resampled texts

Value

a corpus object containing new resampled texts.

Note

Additional resampling units to be added will include fixed length samples and random length samples.

Examples

```
testCorp <- resample(subset(inaugCorpus, Year>2000), 10, "sentences")
testCorpPara <- resample(corpus(uk2010immig), 10, "paragraphs")
names(metadoc(testCorp))
x <- corpus(c("Sentence One C1. Sentence Two C1. Sentence Three C1.",
              "Sentence One C2. Sentence Two C2. Sentence Three C2. Sentence Four C2. Sentence Five C2.",
              docnames=c("docTwo", "docOne")))
testRS <- resample(x, n=3)
metadoc(testRS)
testRS$documents

# tests to see if a corpus contains resampled texts
is.resampled(testCorp)
is.resampled(inaugCorpus)
```

segmentSentence	<i>segment texts into component elements</i>
-----------------	--

Description

Segment text(s) into tokens, sentences, paragraphs, or other sections. `segment` works on a character vector or corpus object, and allows the delimiters to be defined. See details.

Usage

```
segmentSentence(x, delimiter = "[.!?:;]")

segmentParagraph(x, delimiter = "\\n{2}")

segment(x, ...)

## S3 method for class character
segment(x, what = c("tokens", "sentences", "paragraphs",
                    "other"), delimiter = ifelse(what == "tokens", " ", ifelse(what ==
                    "sentences", "[.!?:;]", "\\n{2}")), ...)

## S3 method for class corpus
segment(x, what = c("tokens", "sentences", "paragraphs",
                    "other"), delimiter = ifelse(what == "tokens", " ", ifelse(what ==
                    "sentences", "[.!?:;]", "\\n{2}")), ...)
```

Arguments

<code>x</code>	text or corpus object to be segmented
<code>delimiter</code>	delimiter defined as a regex for segmentation. Each type has its own default, except other, which requires a value to be specified.
<code>...</code>	provides additional arguments to be passed to clean
<code>what</code>	unit of segmentation. Current options are tokens, sentences, paragraphs, and other. Segmenting on other allows segmentation of a text on any user-defined value, and must be accompanied by the <code>delimiter</code> argument.

Details

Tokens are delimited by whitespace. For sentences, the delimiter can be defined by the user. The default for sentences includes ., !, ?, plus ; and :.

For paragraphs, the default is two carriage returns, although this could be changed to a single carriage return by changing the value of `delimiter` to `"\\n{1}"` which is the R version of the [regex](#) for one newline character. (You might need this if the document was created in a word processor, for instance, and the lines were wrapped in the window rather than being hard-wrapped with a newline character.)

Value

`segmentSentence` returns a character vector of sentences that have been segmented

`segmentParagraph` returns a character vector of paragraphs that have been segmented

A list of segmented texts, with each element of the list corresponding to one of the original texts.

Examples

```
# segment sentences of the UK 2010 immigration sections of manifestos
segmentSentence(uk2010immig[1])[1:5] # 1st 5 sentences from first (BNP) text
str(segmentSentence(uk2010immig[1])) # a 143-element char vector
str(segmentSentence(uk2010immig[1:2])) # a 155-element char vector (143+ 12)
# segment paragraphs
segmentParagraph(uk2010immig[3])[1:2] # 1st 2 Paragraphs from 3rd (Con) text
str(segmentParagraph(uk2010immig[3])) # a 12-element char vector
# same as tokenize()
identical(tokenize(uk2010immig, lower=FALSE), segment(uk2010immig, lower=FALSE))

# segment into paragraphs
segment(uk2010immig[3:4], "paragraphs")

# segment a text into sentences
segmentedChar <- segment(uk2010immig, "sentences")
segmentedChar[2]
# segment a corpus into sentences
segmentedCorpus <- segment(corpus(uk2010immig), "sentences")
identical(segmentedCorpus, segmentedChar)
```

settings

Get or set the corpus settings

Description

Get or set the corpus settings

Get or set various settings in the corpus for the treatment of texts, such as rules for stemming, stopwords, collocations, etc.

Get the settings from a which a [dfm](#) was created

Usage

```

settings(x, ...)

## S3 method for class corpus
settings(x, field = NULL, ...)

settings(x, field) <- value

## S3 method for class dfm
settings(x, ...)

```

Arguments

x	object from/to which settings are queried or applied
...	additional arguments
field	string containing the name of the setting to be set or queried settings(x) query the corps settings settings(x, field) <- update the corpus settings for field
value	new setting value

Examples

```

settings(inaugCorpus, "stopwords")
tempdfm <- dfm(inaugCorpus)
tempdfmSW <- dfm(inaugCorpus, stopwords=TRUE)
settings(inaugCorpus, "stopwords") <- TRUE
tempdfmSW <- dfm(inaugCorpus)
tempdfm <- dfm(inaugCorpus, stem=TRUE)
settings(tempdfm)

```

settingsInitialize	<i>settingsInitialize returns a list of legal settings, set to their default values</i>
--------------------	---

Description

settingsInitialize returns a list of legal settings, set to their default values

Usage

```
settingsInitialize()
```

similarity	<i>compute similarities between documents and/or features</i>
------------	---

Description

Compute similarities between documents and/or features from a [dfm](#). Uses the similarity measures defined in [simil](#). See [pr_DB](#) for available distance measures, or how to create your own.

Usage

```
similarity(x, selection, n = 10, margin = c("features", "documents"),
  method = "correlation", sort = TRUE, normalize = TRUE, digits = 4)
```

Arguments

x	a dfm object
selection	character or character vector of document names or feature labels from the dfm
n	the top n most similar items will be returned, sorted in descending order. If n is NULL, return all items.
margin	identifies the margin of the dfm on which similarity will be computed: features for word/term features or documents for documents.
method	a valid method for computing similarity from pr_DB
sort	sort results in descending order if TRUE
normalize	if TRUE, normalize the dfm by term frequency within document (so that the dfm values will be relative term frequency within each document)
digits	digits for rounding results

Value

a named list of the selection labels, with a sorted named vector of similarity measures.

Note

The method for computing feature similarities can be quite slow when there are large numbers of feature types. Future implementations will hopefully speed this up.

Examples

```
# create a dfm from inaugural addresses from Reagan onwards
presDfm <- dfm(subset(inaugCorpus, Year>1980), stopwords=TRUE, stem=TRUE)

# compute some document similarities
similarity(presDfm, "1985-Reagan", n=5, margin="documents")
similarity(presDfm, c("2009-Obama" , "2013-Obama"), n=5, margin="documents")
similarity(presDfm, c("2009-Obama" , "2013-Obama"), n=NULL, margin="documents")
similarity(presDfm, c("2009-Obama" , "2013-Obama"), n=NULL, margin="documents", method="cosine")
similarity(presDfm, "2005-Bush", n=NULL, margin="documents", method="eJaccard", sort=FALSE)

## Not run:
# compute some term similarities
```

```

similarity(presDfm, c("fair", "health", "terror"), method="cosine")

# compare to tm
require(tm)
data("crude")
crude <- tm_map(crude, content_transformer(tolower))
crude <- tm_map(crude, removePunctuation)
crude <- tm_map(crude, removeNumbers)
crude <- tm_map(crude, stemDocument)
tdm <- TermDocumentMatrix(crude)
findAssocs(tdm, c("oil", "opec", "xyz"), c(0.75, 0.82, 0.1))
# in quanteda
crudeDfm <- dfm(corpus(crude))
similarity(crudeDfm, c("oil", "opec", "xyz"), normalize=FALSE, digits=2)

## End(Not run)

```

sort.dfm

sort a dfm by one or more margins

Description

Sorts a [dfm](#) by frequency of total features, total features in documents, or both

Usage

```

## S3 method for class dfm
sort(x, decreasing = TRUE, margin = c("features", "docs",
  "both"), ...)

```

Arguments

x	Document-feature matrix created by dfm
decreasing	TRUE (default) if sort will be in descending order, otherwise sort in increasing order
margin	which margin to sort on features to sort by frequency of features, docs to sort by total feature counts in documents, and both to sort by both
...	additional argumnets passed to base method <code>sort.int</code>

Value

A sorted [dfm](#) matrix object

Author(s)

Ken Benoit

Examples

```
dtm <- dfm(inaugCorpus)
dtm[1:10, 1:5]
dtm <- sort(dtm)
sort(dtm)[1:10, 1:5]
sort(dtm, TRUE, "both")[1:10, 1:5] # note that the decreasing=TRUE argument
                                   # must be second, because of the order of the
                                   # formals in the generic method of sort()
```

statLexdiv	<i>calculate lexical diversity</i>
------------	------------------------------------

Description

Calculate the lexical diversity or complexity of text(s).

Usage

```
statLexdiv(x, ...)

## S3 method for class dfm
statLexdiv(x, measure = c("TTR", "C", "R", "CTTR", "U", "S",
  "Maas"), log.base = 10, ...)

## S3 method for class numeric
statLexdiv(x, measure = c("TTR", "C", "R", "CTTR", "U", "S",
  "Maas"), log.base = 10, ...)
```

Arguments

x	a document-feature matrix object
...	additional arguments
measure	A character vector defining the measure to calculate.
log.base	A numeric value defining the base of the logarithm.

Details

statLexdiv calculates a variety of proposed indices for lexical diversity. In the following formulae, N refers to the total number of tokens, and V to the number of types:

"TTR": The ordinary *Type-Token Ratio*:

$$TTR = \frac{V}{N}$$

"C": Herdan's C (Herdan, 1960, as cited in Tweedie & Baayen, 1998; sometimes referred to as *LogTTR*):

$$C = \frac{\lg V}{\lg N}$$

"R": Guiraud's *Root TTR* (Guiraud, 1954, as cited in Tweedie & Baayen, 1998):

$$R = \frac{V}{\sqrt{N}}$$

"CTTR": Carroll's *Corrected TTR*:

$$CTTR = \frac{V}{\sqrt{2N}}$$

"U": Dugast's *Uber Index* (Dugast, 1978, as cited in Tweedie & Baayen, 1998):

$$U = \frac{(\lg N)^2}{\lg N - \lg V}$$

"S": Summer's index:

$$S = \frac{\lg \lg V}{\lg \lg N}$$

"K": Yule's *K* (Yule, 1944, as cited in Tweedie & Baayen, 1998) is calculated by:

$$K = 10^4 \times \frac{(\sum_{X=1}^X f_X X^2) - N}{N^2}$$

where N is the number of tokens, X is a vector with the frequencies of each type, and f_X is the frequencies for each X .

"Maas": Maas' indices (a , $\lg V_0$ & $\lg_e V_0$):

$$a^2 = \frac{\lg N - \lg V}{\lg N^2}$$

$$\lg V_0 = \frac{\lg V}{\sqrt{1 - \frac{\lg V^2}{\lg N^2}}}$$

The measure was derived from a formula by Muller (1969, as cited in Maas, 1972). $\lg_e V_0$ is equivalent to $\lg V_0$, only with e as the base for the logarithms. Also calculated are a , $\lg V_0$ (both not the same as before) and V' as measures of relative vocabulary growth while the text progresses. To calculate these measures, the first half of the text and the full text will be examined (see Maas, 1972, p. 67 ff. for details). Note: for the current method (for a dfm) there is no computation on separate halves of the text.

Value

a vector of lexical diversity statistics, each corresponding to an input document

Note

This implements only the static measures of lexical diversity, not more complex measures based on windows of text such as the Mean Segmental Type-Token Ratio, the Moving-Average Type-Token Ratio (Covington & McFall, 2010), the MLTD or MLTD-MA (Moving-Average Measure of Textual Lexical Diversity) proposed by McCarthy & Jarvis (2010) or Jarvis (no year), or the HD-D version of vocd-D (see McCarthy & Jarvis, 2007). These are available from the package **korRpus**.

Author(s)

Kenneth Benoit, adapted from the S4 class implementation written by Meik Michalke in the **korRpus** package.

References

Covington, M.A. & McFall, J.D. (2010). Cutting the Gordian Knot: The Moving-Average Type-Token Ratio (MATTR). *Journal of Quantitative Linguistics*, 17(2), 94–100.

Maas, H.-D., (1972). \ "Uber den Zusammenhang zwischen Wortschatzumfang und L\ "ange eines Textes. *Zeitschrift f\ "ur Literaturwissenschaft und Linguistik*, 2(8), 73–96.

McCarthy, P.M. & Jarvis, S. (2007). vocd: A theoretical and empirical evaluation. *Language Testing*, 24(4), 459–488.

McCarthy, P.M. & Jarvis, S. (2010). MTLT, vocd-D, and HD-D: A validation study of sophisticated approaches to lexical diversity assessment. *Behaviour Research Methods*, 42(2), 381–392.

Michalke, Meik. (2014) *koRpus: An R Package for Text Analysis*. Version 0.05-5. <http://reaktanz.de/?c=hacking&s=koRpus>

Tweedie, F.J. & Baayen, R.H. (1998). How Variable May a Constant Be? Measures of Lexical Richness in Perspective. *Computers and the Humanities*, 32(5), 323–352.

Examples

```
mydfm <- dfm(inaugCorpus)
mydfmSW <- dfm(inaugCorpus, stopwords=TRUE)
results <- data.frame(TTR = statLexdiv(mydfm, "TTR"),
                      CTTR = statLexdiv(mydfm, "CTTR"),
                      U = statLexdiv(mydfm, "U"),
                      TTRs = statLexdiv(mydfmSW, "TTR"),
                      CTTRs = statLexdiv(mydfmSW, "CTTR"),
                      Us = statLexdiv(mydfmSW, "U"))

results
cor(results)
t(statLexdiv(mydfmSW, "Maas"))
```

stopwords	A named list containing common stopwords in 14 languages
-----------	--

Description

SMART English stopwords from the SMART information retrieval system (obtained from <http://jmlr.csail.mit.edu/papers/smart-stop-list/english.stop>) and a set of stopword lists from the Snowball stemmer project in different languages (obtained from http://svn.tartarus.org/snowball/trunk/website/algorithms/*/stop.txt). Supported languages are danish, dutch, english, finnish, french, german, hungarian, italian, norwegian, portuguese, russian, spanish, and swedish. Language names are case sensitive. Alternatively, their IETF language tags may be used.

stopwordsGet	<i>access stopwords</i>
--------------	-------------------------

Description

This function retrieves stopwords from the type specified in the `kind` argument and returns the stopword list as a character vector. The default is English. See [stopwords](#) for information about the list.

Usage

```
stopwordsGet(kind = "english")
```

Arguments

<code>kind</code>	The pre-set kind of stopwords (as a character string)
-------------------	---

Value

a character vector or dfm with stopwords removed

Examples

```
stopwordsGet()
stopwordsGet("italian")
```

stopwordsRemove	<i>remove stopwords from a text or dfm</i>
-----------------	--

Description

This function takes a character vector or [dfm](#) and removes words in the remove common or 'semantically empty' words from a text. See [stopwordsGet](#) for the information about the default lists.

Usage

```
stopwordsRemove(text, stopwords = NULL)

## S3 method for class character
stopwordsRemove(text, stopwords = NULL)

## S3 method for class dfm
stopwordsRemove(text, stopwords = NULL)
```

Arguments

<code>text</code>	Text from which stopwords will be removed
<code>stopwords</code>	Character vector of stopwords to remove - if none is supplied, a default set of English stopwords is used

Details

This function takes a character vector 'text' and removes words in the list provided in stopwords. If no list of stopwords is provided a default list for English is used. The function [stopwordsGet](#) can load a default set of stopwords for many languages.

Value

a character vector or dfm with stopwords removed

Examples

```
## examples for character objects
someText <- "Here is an example of text containing some stopwords we want to remove."
itText <- "Ecco un esempio di testo contenente alcune parole non significative che vogliamo rimuovere."
stopwordsRemove(someText)
stopwordsRemove(someText, stopwordsGet("SMART"))
stopwordsRemove(itText, stopwordsGet("italian"))
stopwordsRemove(someText, c("containing", "example"))

## example for dfm objects
docmat <- dfm(uk2010immig)
docmatNostopwords <- stopwordsRemove(docmat)
dim(docmat)
dim(docmatNostopwords)
dim(stopwordsRemove(docmat, stopwordsGet("SMART")))
```

subset.corpus	<i>extract a subset of a corpus</i>
---------------	-------------------------------------

Description

Works just like the normal subset command but for corpus objects

Usage

```
## S3 method for class corpus
subset(x, subset = NULL, select = NULL, ...)
```

Arguments

x	corpus object to be subsetted.
subset	logical expression indicating elements or rows to keep: missing values are taken as false.
select	expression, indicating the attributes to select from the corpus
...	additional arguments affecting the summary produced

Value

corpus object

Examples

```
summary(subset(inaugCorpus, Year>1980))
summary(subset(inaugCorpus, Year>1930 & President=="Roosevelt", select=Year))
```

summary.corpus

*Corpus summary***Description**

Displays information about a corpus object, including attributes and metadata such as date of number of texts, creation and source.

Usage

```
## S3 method for class corpus
summary(object, n = 100, verbose = TRUE,
        showmeta = FALSE, ...)
```

Arguments

object	corpus to be summarized
n	maximum number of texts to describe, default=100
verbose	FALSE to turn off printed output
showmeta	TRUE to include document-level meta-data
...	additional arguments affecting the summary produced

Examples

```
summary(inaugCorpus)
summary(inaugCorpus, n=10)
mycorpus <- corpus(uk2010immig, docvars=data.frame(party=names(uk2010immig)), enc="UTF-8")
summary(mycorpus, showmeta=TRUE) # show the meta-data
mysummary <- summary(mycorpus, verbose=FALSE) # (quietly) assign the results
mysummary$Types / mysummary$Tokens           # crude type-token ratio
```

syllableCounts

*A named list mapping words to counts of their syllables***Description**

A named list mapping words to counts of their syllables, generated from the CMU pronunciation dictionary

References

<http://www.speech.cs.cmu.edu/cgi-bin/cmudict>

Examples

```
data(syllableCounts)
syllableCounts["sixths"]
syllableCounts["onomatopoeia"]
```

texts	<i>get or set corpus texts</i>
-------	--------------------------------

Description

Get or replace the texts in a quanteda corpus object.

Usage

```
texts(corp)

texts(corp) <- value
```

Arguments

corp	A quanteda corpus object
value	character vector of the new texts

Value

For texts, a character vector of the texts in the corpus.
 For texts <-, the corpus with the updated texts.

Examples

```
texts(inaugCorpus)[1]
sapply(texts(inaugCorpus), nchar) # length in characters of the inaugural corpus texts

## this doesnt work yet - need to overload [ for this replacement function
# texts(inaugTexts)[55] <- "GW Bushs second inaugural address, the condensed version."
```

tf	<i>normalizes the term frequencies a dfm</i>
----	--

Description

Returns a matrix of term weights, as a [dfm](#) object

Usage

```
tf(x)
```

Arguments

x	Document-feature matrix created by dfm
---	--

Value

A dfm matrix object where values are relative term proportions within the document

Author(s)

Ken Benoit

Examples

```
data(inaugCorpus)
dtm <- dfm(inaugCorpus)
dtm[1:10, 100:110]
tf(dtm)[1:10, 100:110]
```

tfidf	<i>compute the tf-idf weights of a dfm</i>
-------	--

Description

Returns a matrix of tf-idf weights, as a [dfm](#) object

Usage

```
tfidf(x, normalize = TRUE)

## S3 method for class dfm
tfidf(x, normalize = TRUE)
```

Arguments

x	document-feature matrix created by dfm
normalize	whether to normalize term frequency by document totals

Value

A dfm matrix object where values are tf-idf weights

Author(s)

Ken Benoit

Examples

```
data(inaugCorpus)
dtm <- dfm(inaugCorpus)
dtm[1:10, 100:110]
tfidf(dtm)[1:10, 100:110]
tfidf(dtm, normalize=FALSE)[1:10, 100:110]
```

tokenize	<i>tokenize a set of texts</i>
----------	--------------------------------

Description

Tokenize the texts from a character vector or from a corpus.

Usage

```
tokenize(x, ...)

## S3 method for class character
tokenize(x, simplify = FALSE, sep = " ", ...)

## S3 method for class corpus
tokenize(x, ...)
```

Arguments

x	The text(s) or corpus to be tokenized
...	additional arguments passed to clean
simplify	If TRUE, return a character vector of tokens rather than a list of length ndoc (texts), with each element of the list containing a character vector of the tokens corresponding to that text.
sep	by default, tokenize expects a ‘white-space’ delimiter between tokens. Alternatively, sep can be used to specify another character which delimits fields.

Value

A list of length [ndoc](#)(x) of the tokens found in each text.

A list of length [ndoc](#)(texts) of the tokens found in each text.

Examples

```
# same for character vectors and for lists
tokensFromChar <- tokenize(inaugTexts)
tokensFromCorp <- tokenize(inaugCorpus)
identical(tokensFromChar, tokensFromCorp)
str(tokensFromChar)
# returned as a list
head(tokenize(inaugTexts[57])[[1]], 10)
# returned as a character vector using simplify=TRUE
head(tokenize(inaugTexts[57], simplify=TRUE), 10)

# demonstrate some options with clean
head(tokenize(inaugTexts[57], simplify=TRUE, lower=FALSE), 30)
```

topfeatures	<i>list the most frequent features</i>
-------------	--

Description

List the most frequently occurring features in a [dfm](#)

Usage

```
topfeatures(x, n = 10, decreasing = TRUE, ci = 0.95)
```

```
## S3 method for class dfm
```

```
topfeatures(x, n = 10, decreasing = TRUE, ci = 0.95)
```

Arguments

x	the object whose features will be returned
n	how many top features should be returned
decreasing	If TRUE, return the n most frequent features, if FALSE, return the n least frequent features
ci	confidence interval from 0-1.0 for use if dfm is resampled

Value

A named numeric vector of feature counts, where the names are the feature labels.

Examples

```
topfeatures(dfm(inaugCorpus))
topfeatures(dfm(inaugCorpus, stopwords=TRUE))
# least frequent features
topfeatures(dfm(inaugCorpus), decreasing=FALSE)
```

trimdfm	<i>Trim a dfm based on a subset of features and words</i>
---------	---

Description

Returns a document by feature matrix reduced in size based on document and term frequency, and/or subsampling.

Usage

```
trimdfm(x, minCount = 5, minDoc = 5, sample = NULL, verbose = TRUE)
```

Arguments

<code>x</code>	document-feature matrix created by dfm
<code>minCount</code>	minimum feature count
<code>minDoc</code>	minimum number of documents in which a feature appears
<code>sample</code>	how many features to retain (based on random selection)
<code>verbose</code>	print messages

Value

A [dfm](#) object reduced in size.

Author(s)

Ken Benoit adapted from code by Will Lowe (see [trim](#))

Examples

```
data(inaugCorpus)
dtm <- dfm(inaugCorpus)
dim(dtm)
dtmReduced <- trimdfm(dtm, minCount=10, minDoc=2) # only words occurring at least 5 times and in at least 2
dim(dtmReduced)
dtmSampled <- trimdfm(dtm, sample=200) # top 200 words
dim(dtmSampled) # 196 x 200 words
```

uk2010immig

Immigration-related sections of 2010 UK party manifestos

Description

Extracts from the election manifestos of 9 UK political parties from 2010, related to immigration or asylum-seekers.

Format

A named character vector of plain ASCII texts

Examples

```
data(uk2010immig)
uk2010immigCorpus <- corpus(uk2010immig, docvars=list(party=names(uk2010immig)))
language(uk2010immigCorpus) <- "english"
encoding(uk2010immigCorpus) <- "UTF-8"
summary(uk2010immigCorpus)
```

wordstem	<i>stem words</i>
----------	-------------------

Description

Apply a stemmer to words. This is a wrapper to [wordStem](#) designed to allow this function to be called without loading the entire **SnowballC** package. [wordStem](#) uses Dr. Martin Porter's stemming algorithm and the C libstemmer library generated by Snowball.

Usage

```
wordstem(words, language = "porter")
```

Arguments

words	a character vector of words whose stems are to be extracted.
language	the name of a recognized language, as returned by <code>getStemLanguages</code> , or a two- or three-letter ISO-639 code corresponding to one of these languages (see references for the list of codes)

Value

A character vector with as many elements as there are in the input vector with the corresponding elements being the stem of the word. Elements of the vector are converted to UTF-8 encoding before the stemming is performed, and the returned elements are marked as such when they contain non-ASCII characters.

See Also

[wordStem](#); <http://snowball.tartarus.org/>.

Examples

```
# Simple example
wordstem(c("win", "winning", "winner"))
```

zipfiles	<i>unzip a zipped collection of text files and return the directory</i>
----------	---

Description

zipfiles extracts a set of text files in a zip archives, and returns the name of the temporary directory where they are stored. It can be passed to [corpus.directory](#) for import.

Usage

```
zipfiles(zfile = NULL, ...)
```

Arguments

zfile	a character string specifying the name (including path) of the zipped file, or a URL naming the file (see example); or NULL to use a GUI to choose a file from disk
...	additional arguments passed to unzip

Value

a [directory](#) class object containing the unzipped files

Examples

```
## Not run:
# from a zip file on the web
myzipcorp <- corpus(zipfiles("http://kenbenoit.net/files/EUcoalsubsidies.zip"),
  notes="From some EP debate about coal mine subsidies")
docvars(myzipcorp, "speakername") <- docnames(myzipcorp)
summary(myzipcorp)

# call up interactive user input
myzipcorp <- corpus(zipfiles())

## End(Not run)
```


Index

- `+.corpus (corpus)`, 7
- `as.dfm (dfm)`, 11
- `as.DocumentTermMatrix`, 13, 14
- `attributes`, 12
- `bigrams`, 3
- `changeunits`, 4
- `clean`, 4, 11, 31, 44
- `collocations`, 5
- `corpus`, 7, 7
- `corpus.directory`, 47
- `countSyllables`, 10
- `data.frame`, 8
- `describeTexts`, 10
- `dfm`, 8, 11, 11, 12–14, 17, 27, 32, 34, 35, 39, 42, 43, 45, 46
- `dfm2ldaformat`, 13
- `dfm2tmformat`, 13
- `directory`, 7, 14, 48
- `docnames`, 14
- `docnames<- (docnames)`, 14
- `document-feature matrix object`, 36
- `DocumentTermMatrix`, 13
- `docvars`, 8, 9, 15
- `docvars<- (docvars)`, 15
- `Encoding`, 16, 19, 20
- `encoding`, 9, 16
- `encoding<- (encoding)`, 16
- `features (features.dfm)`, 17
- `features.dfm`, 17
- `file`, 14
- `flatten.dictionary`, 17
- `getRootFileNames`, 18
- `getTextDir`, 19
- `getTextDirGui`, 19
- `getTextFiles`, 20
- `getTweets`, 21
- `iconv`, 8
- `iconvlist`, 8, 16, 28, 29
- `inaugCorpus`, 21
- `inaugTexts (inaugCorpus)`, 21
- `is.corpus (corpus)`, 7
- `is.dfm (dfm)`, 11
- `is.resampled (resample)`, 30
- `kwic`, 22
- `language`, 9, 23
- `language<- (language)`, 23
- `lda.collapsed.gibbs.sampler`, 13
- `metacorporus`, 9, 23
- `metacorporus<- (metacorporus)`, 23
- `metadoc`, 8, 9, 16, 23, 24
- `metadoc<- (metadoc)`, 24
- `ndoc`, 25, 44
- `ngrams`, 25
- `nresample`, 26
- `plot.dfm`, 27
- `pr_DB`, 34
- `print.dfm`, 27
- `quanteda`, 21, 28
- `quanteda-package (quanteda)`, 28
- `readLIWCdict`, 28
- `readWStatDict`, 29
- `regex`, 31, 32
- `regular expression`, 5, 8, 19
- `resample`, 12, 30
- `segment`, 4, 30
- `segment (segmentSentence)`, 31
- `segmentParagraph (segmentSentence)`, 31
- `segmentSentence`, 31
- `settings`, 8, 9, 12, 28, 32
- `settings<- (settings)`, 32
- `settingsInitialize`, 33
- `simil`, 34
- `similarity`, 34
- `simple triplet matrix`, 13

sort.dfm, [35](#)
statLexdiv, [36](#)
stopwords, [11](#), [38](#), [39](#)
stopwordsGet, [39](#), [39](#), [40](#)
stopwordsRemove, [39](#)
strheight, [27](#)
strwidth, [27](#)
subset.corpus, [40](#)
summary.corpus, [41](#)
syllableCounts, [41](#)

text, [27](#)
texts, [8](#), [9](#), [42](#)
texts<- (texts), [42](#)
tf, [42](#)
tfidf, [43](#)
tokenise (tokenize), [44](#)
tokenize, [3](#), [26](#), [44](#)
topfeatures, [45](#)
trim, [46](#)
trimdfm, [45](#)

uk2010immig, [46](#)
unlist, [23](#)
unzip, [48](#)

VCorpus, [7](#), [9](#)

wordcloud, [27](#)
wordStem, [47](#)
wordstem, [47](#)

zipfiles, [7](#), [47](#)